

Building and deploying microservices with event sourcing, CQRS and Docker

Chris Richardson

Author of POJOs in Action

Founder of the original CloudFoundry.com

Twitter: @crichtardson

chris@chrisrichardson.net

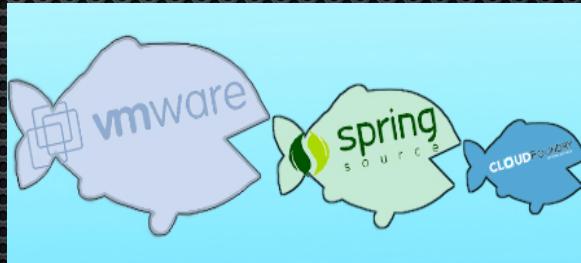
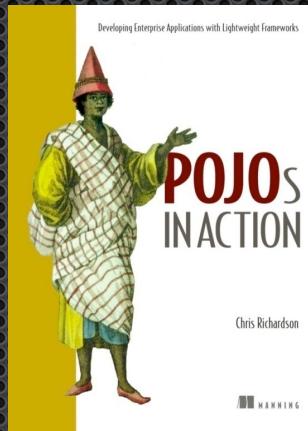
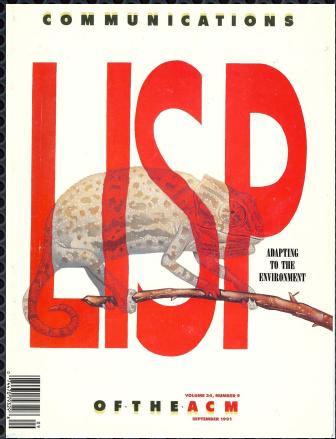
<http://plainoldobjects.com>

<http://microservices.io>

Presentation goal

Share my experiences with building and deploying an application using Scala, functional domain models, microservices, event sourcing, CQRS, and Docker

About Chris



@crichtson

About Chris

- Founder of a buzzword compliant (stealthy, social, mobile, big data, machine learning, ...) startup
- Consultant helping organizations improve how they architect and deploy applications using cloud, micro services, polyglot applications, NoSQL, ...
- Creator of <http://microservices.io>

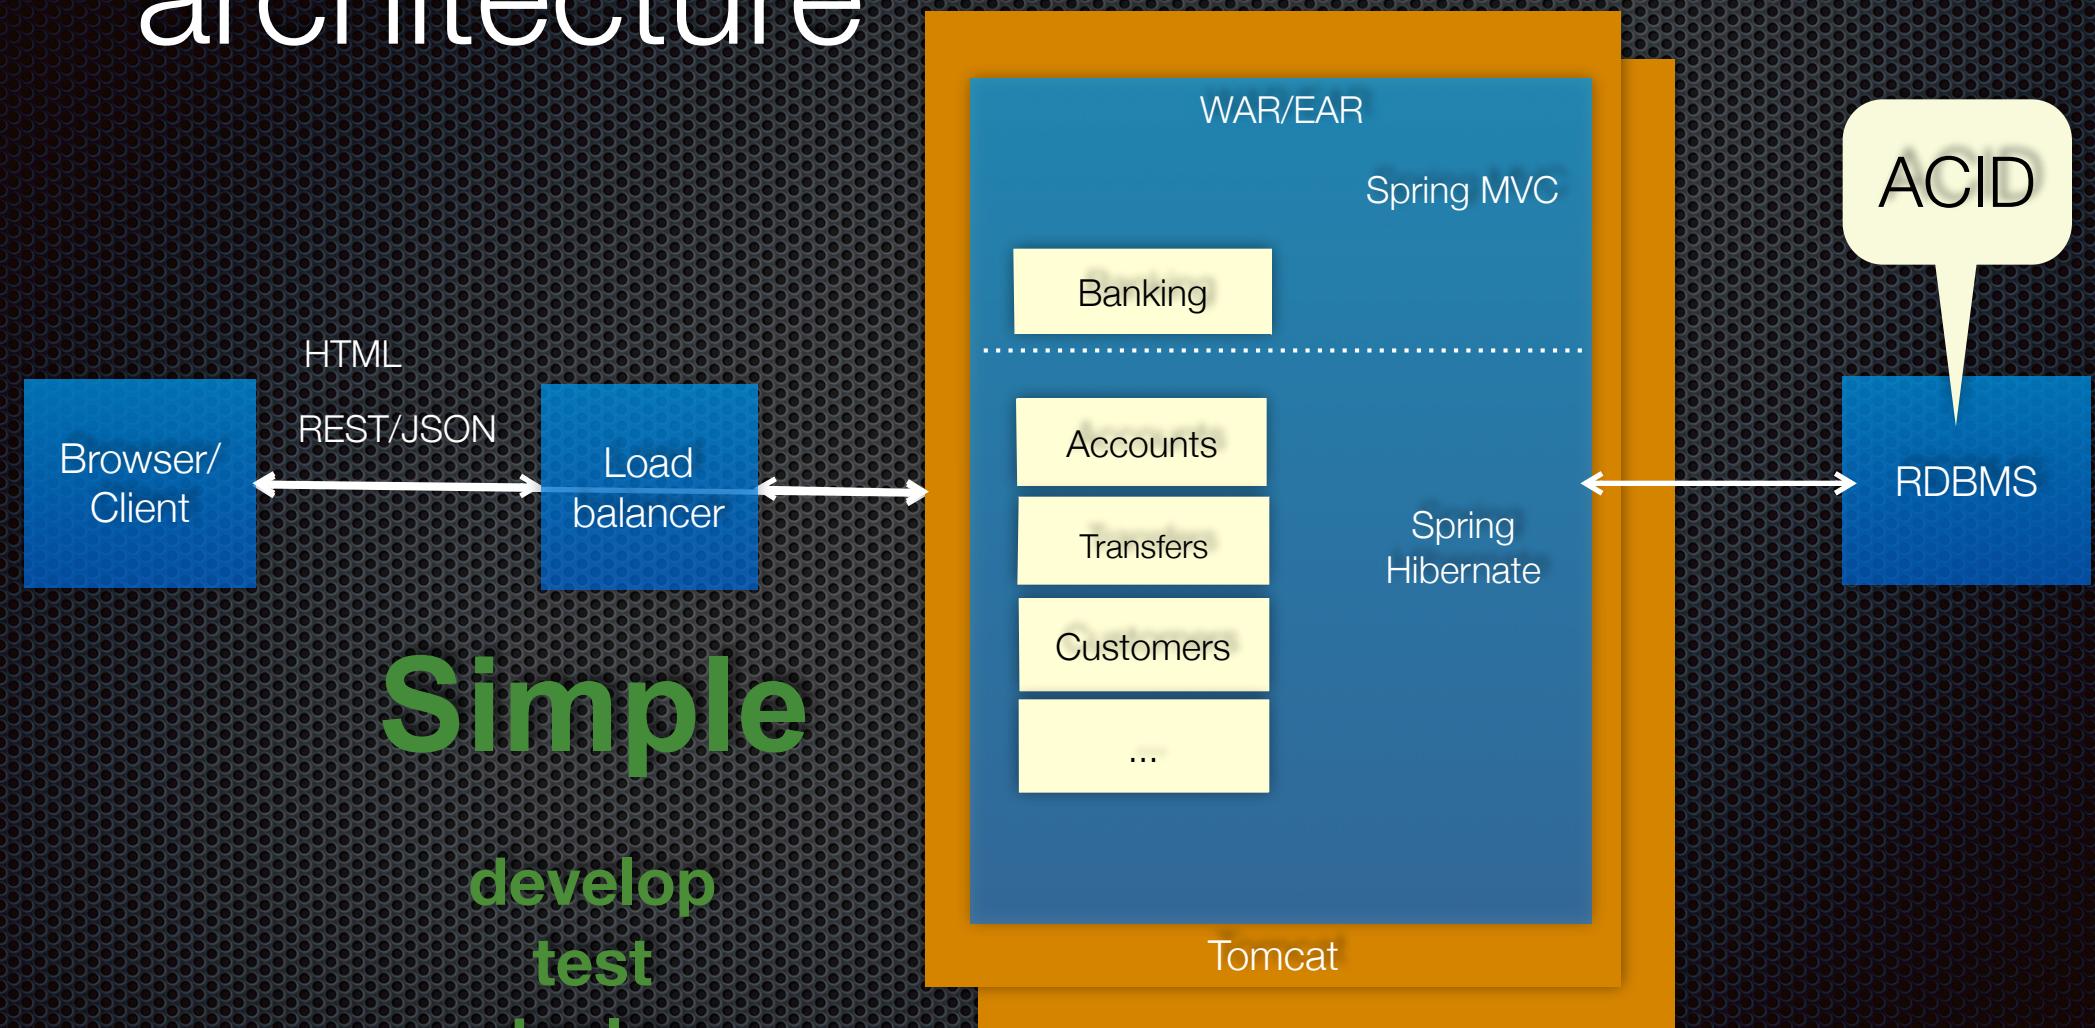
For more information

- <https://github.com/cer/event-sourcing-examples>
- <http://microservices.io>
- <http://plainoldobjects.com/>
- <https://twitter.com/crichardson>

Agenda

- Why build event-driven microservices?
- Overview of event sourcing
- Designing microservices with event sourcing
- Implementing queries in an event sourced application
- Building and deploying microservices

Traditional application architecture



Simple
develop
test
deploy
scale

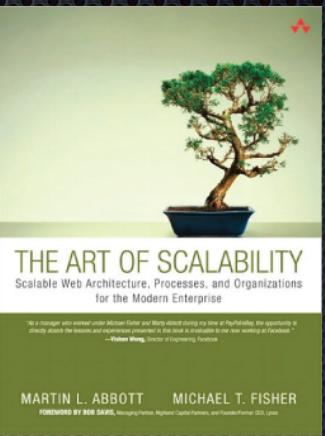
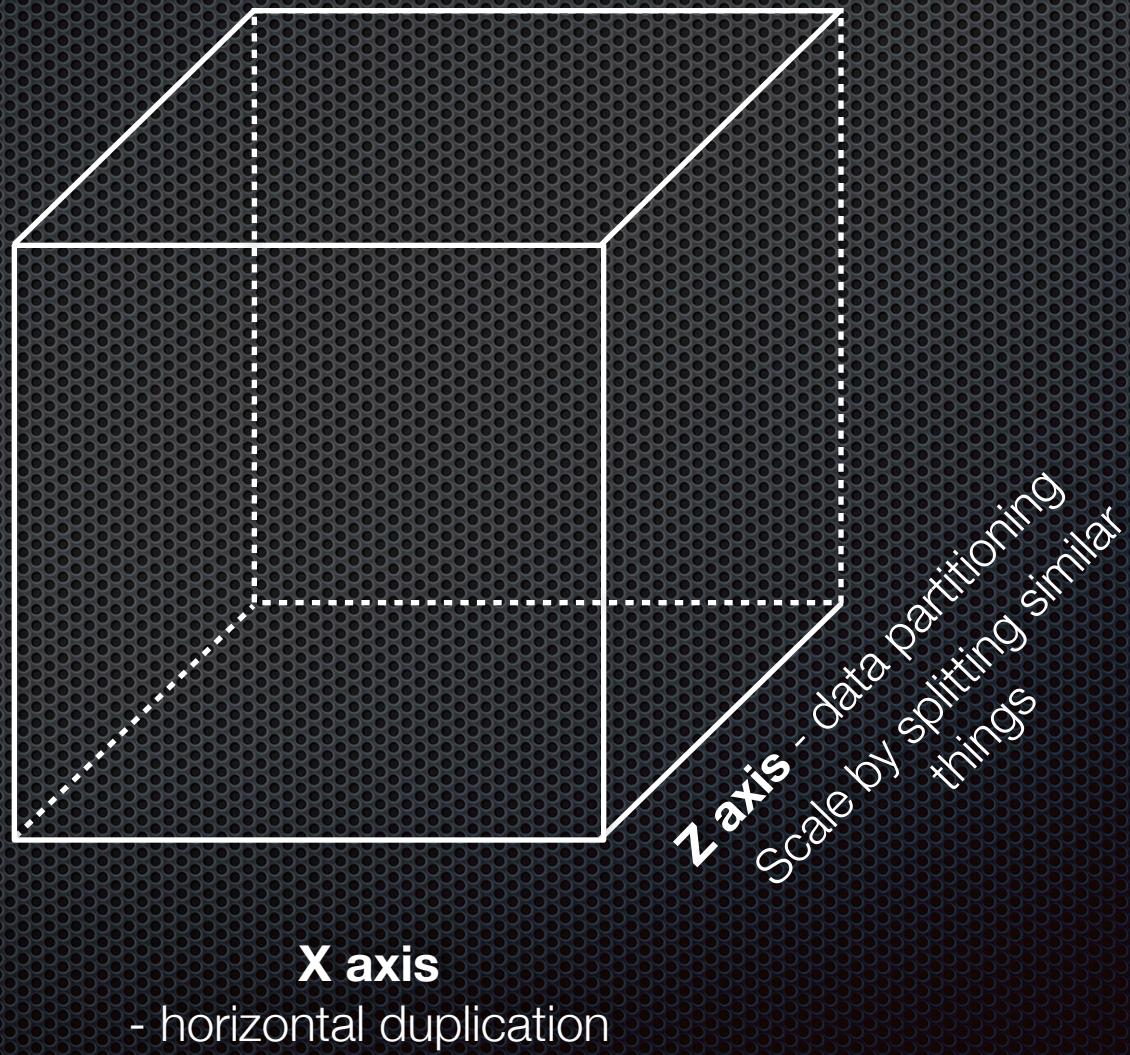
Limitations of the monolithic architecture

- Intimidates developers
- Obstacle to frequent deployments
- Overloads your IDE and container
- Obstacle to scaling development
- Modules having conflicting scaling requirements
- Requires long-term commitment to a technology stack

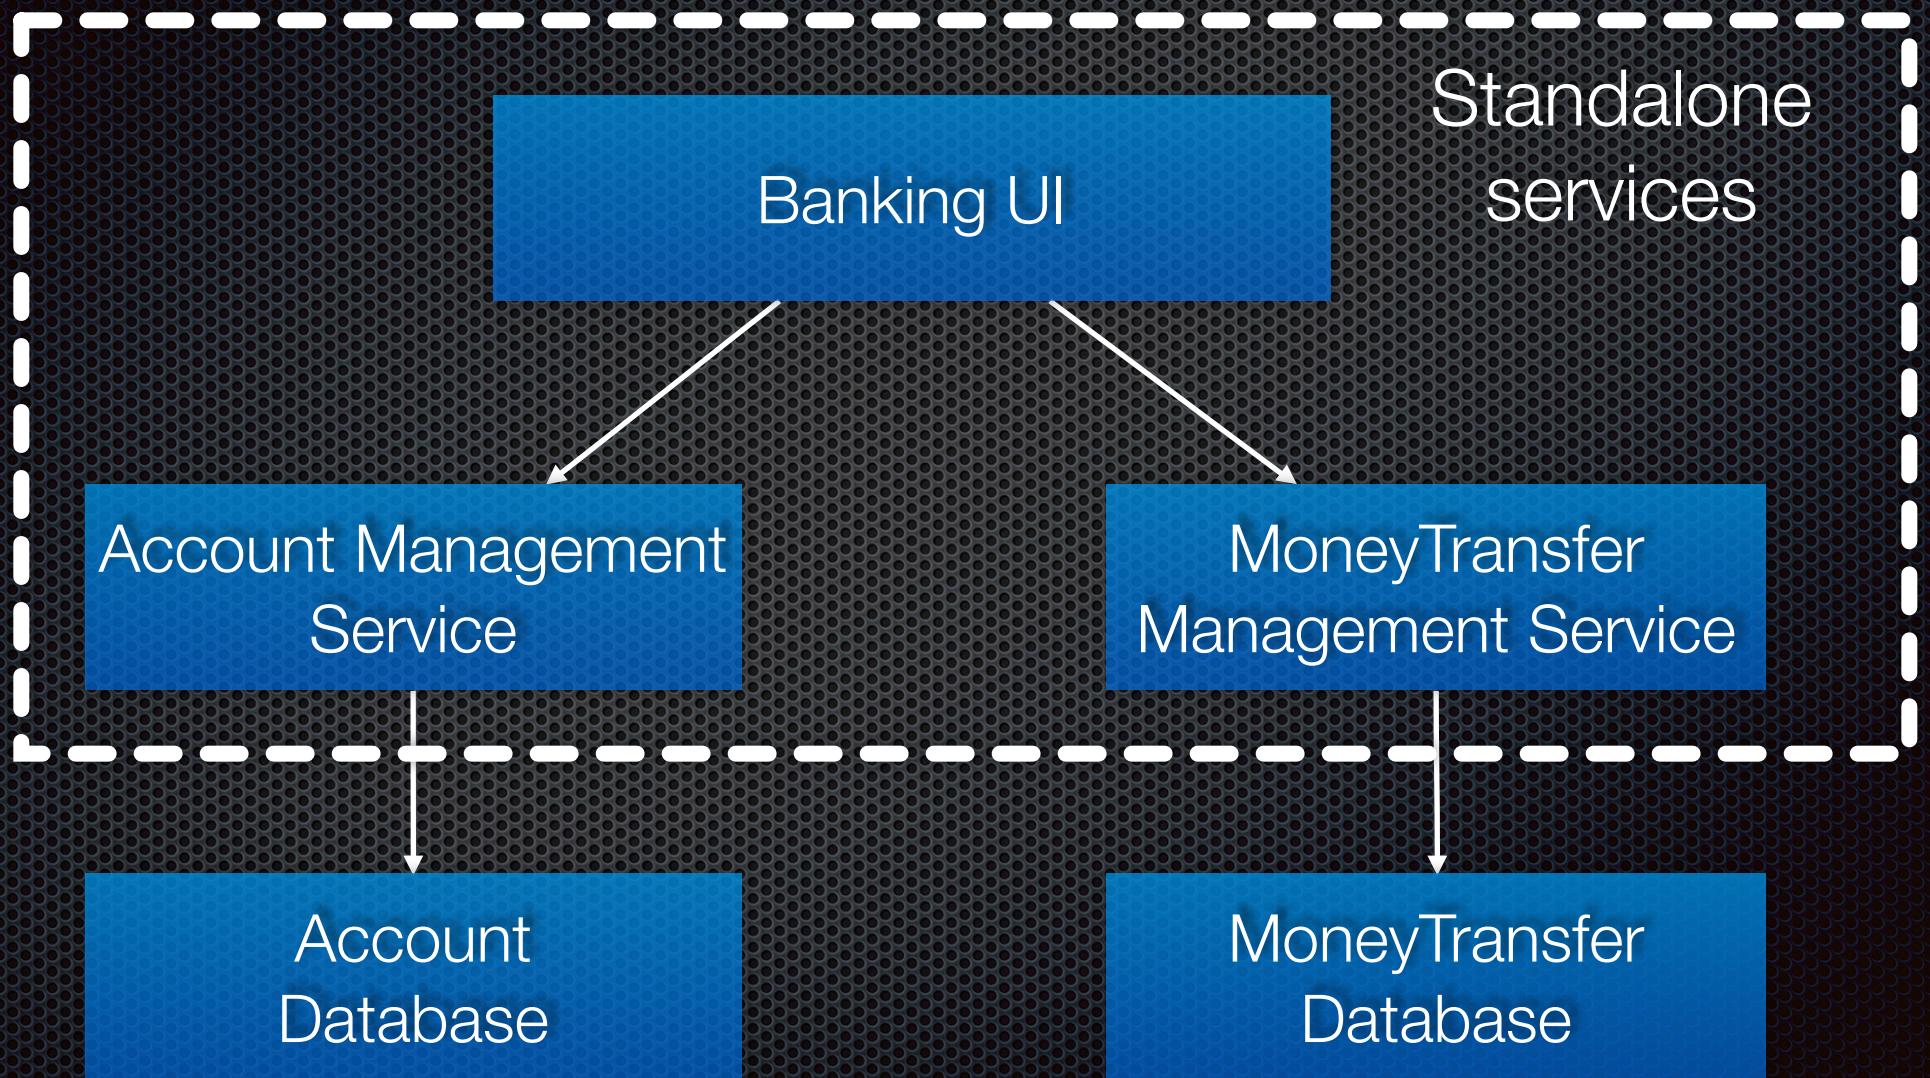
Apply the scale cube

Y axis -
functional
decomposition

Scale by
splitting
different things



Use a microservice architecture



Limitations of a single relational database

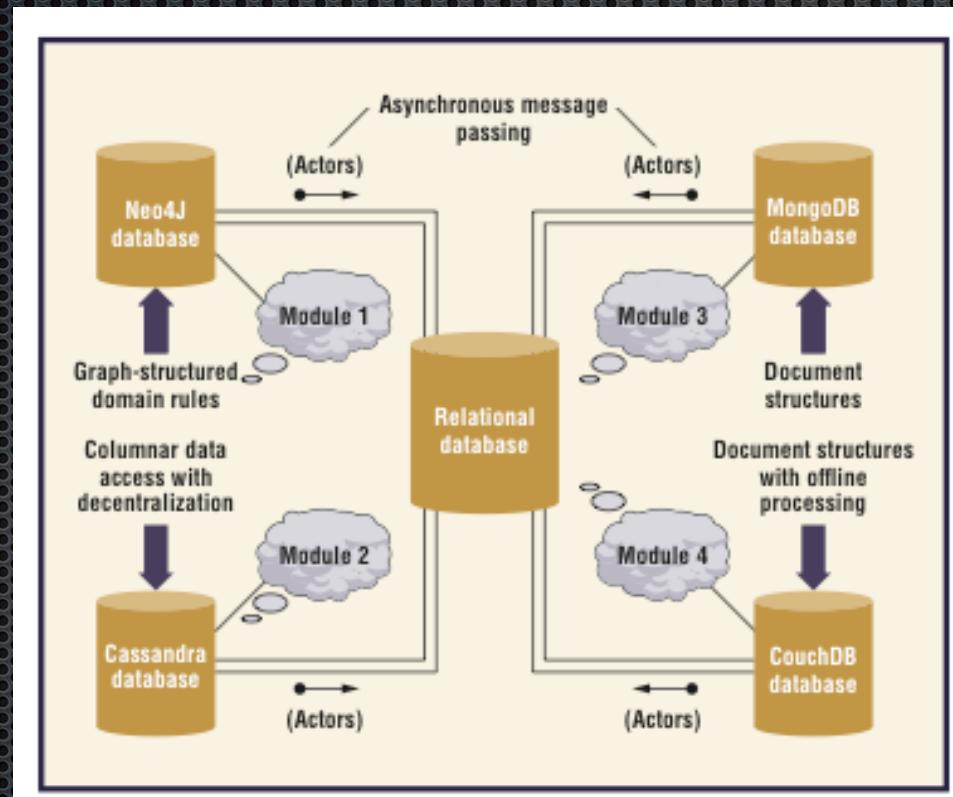
- Scalability
- Distribution
- Schema updates
- O/R impedance mismatch
- Handling semi-structured data

Use a sharded relational database

Use NoSQL databases

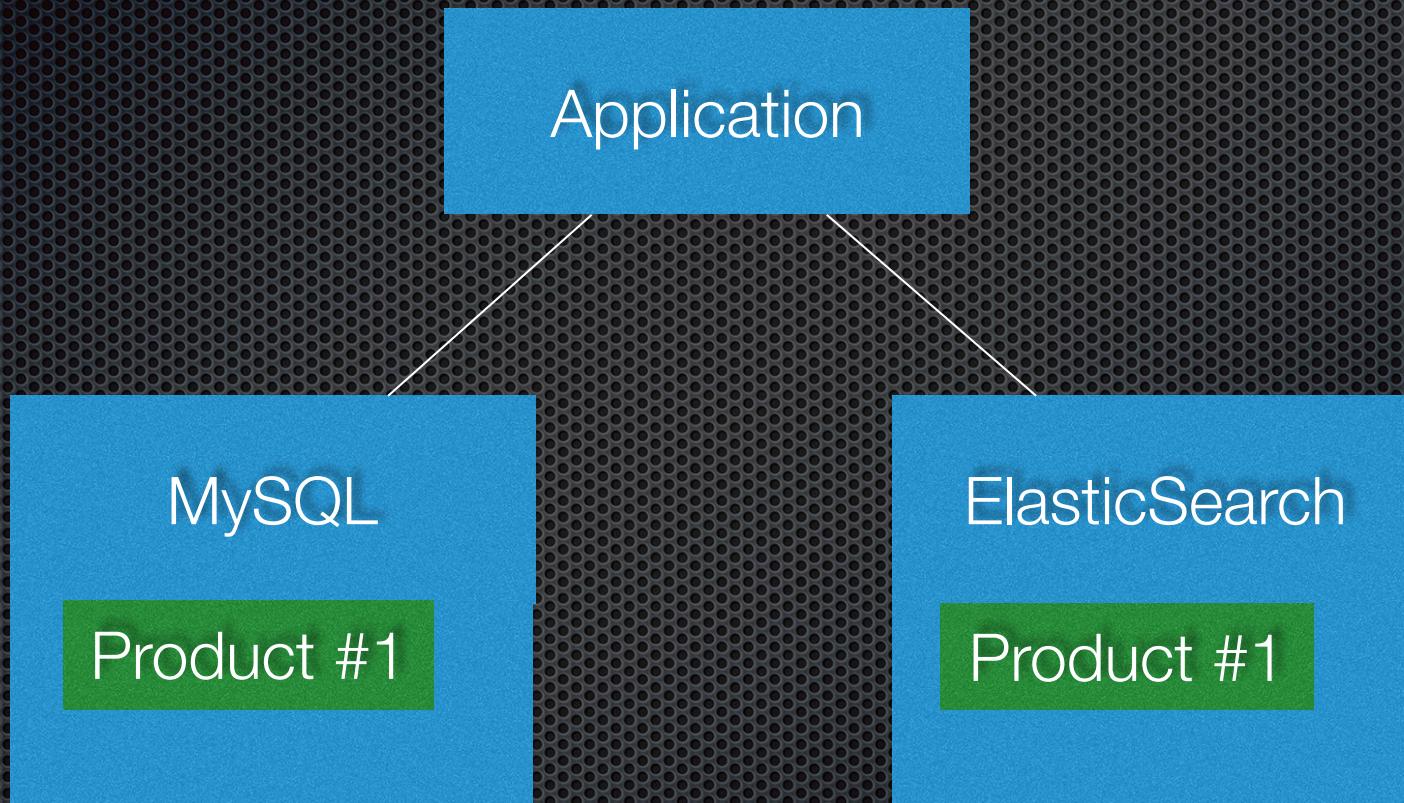
- Avoids the limitations of RDBMS
- For example,
 - text search ⇒ Solr/Cloud Search
 - social (graph) data ⇒ Neo4J
 - highly distributed/available database ⇒ Cassandra
 - ...

Different modules use different types of databases



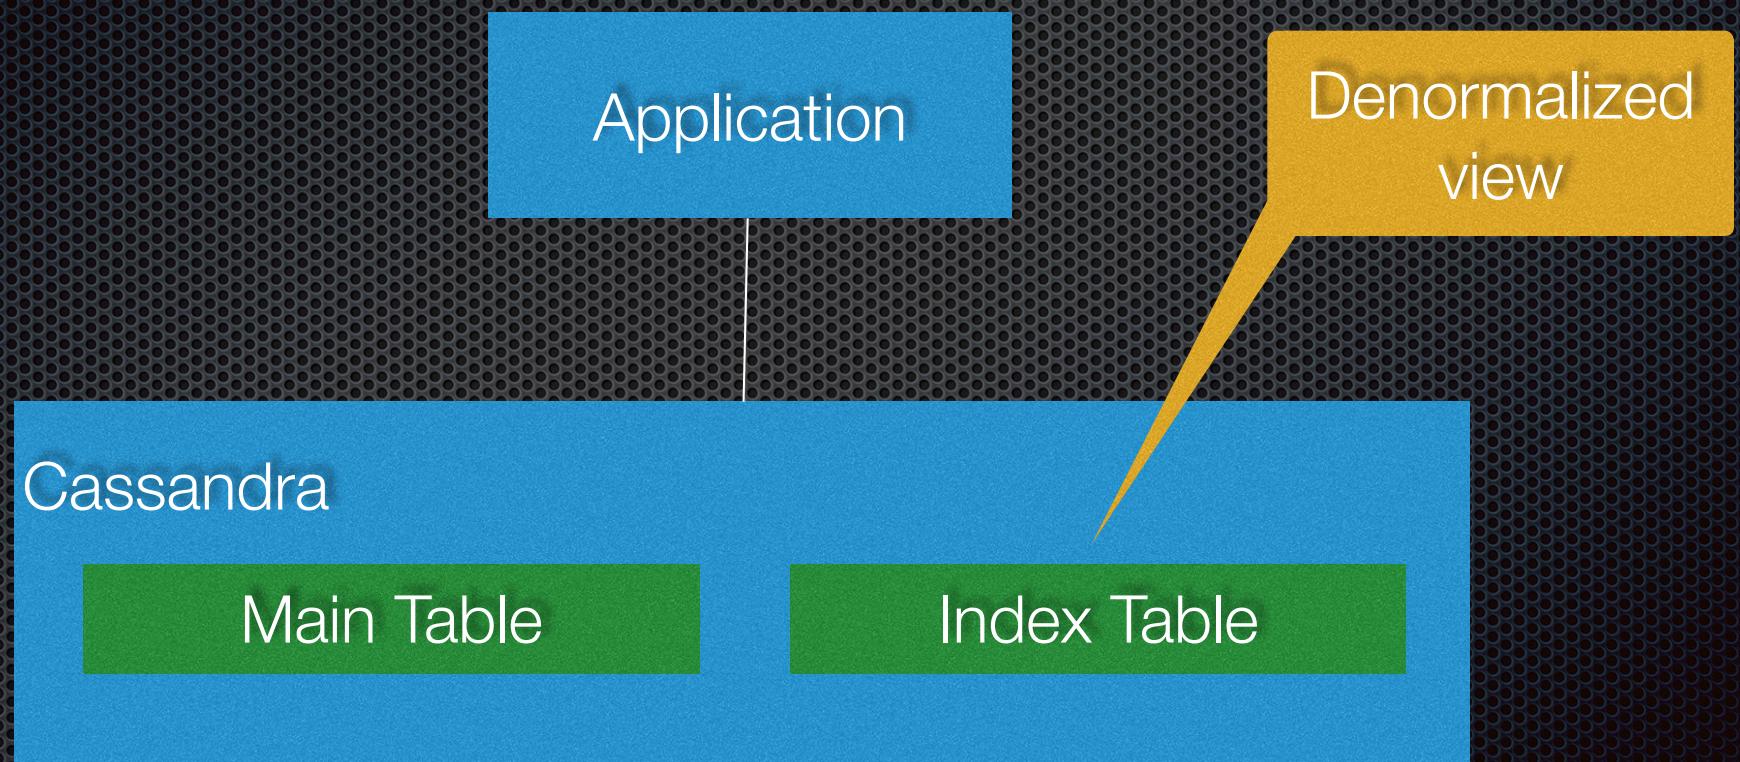
But this results in distributed
data management problems

Example #1 - SQL + Text Search engine



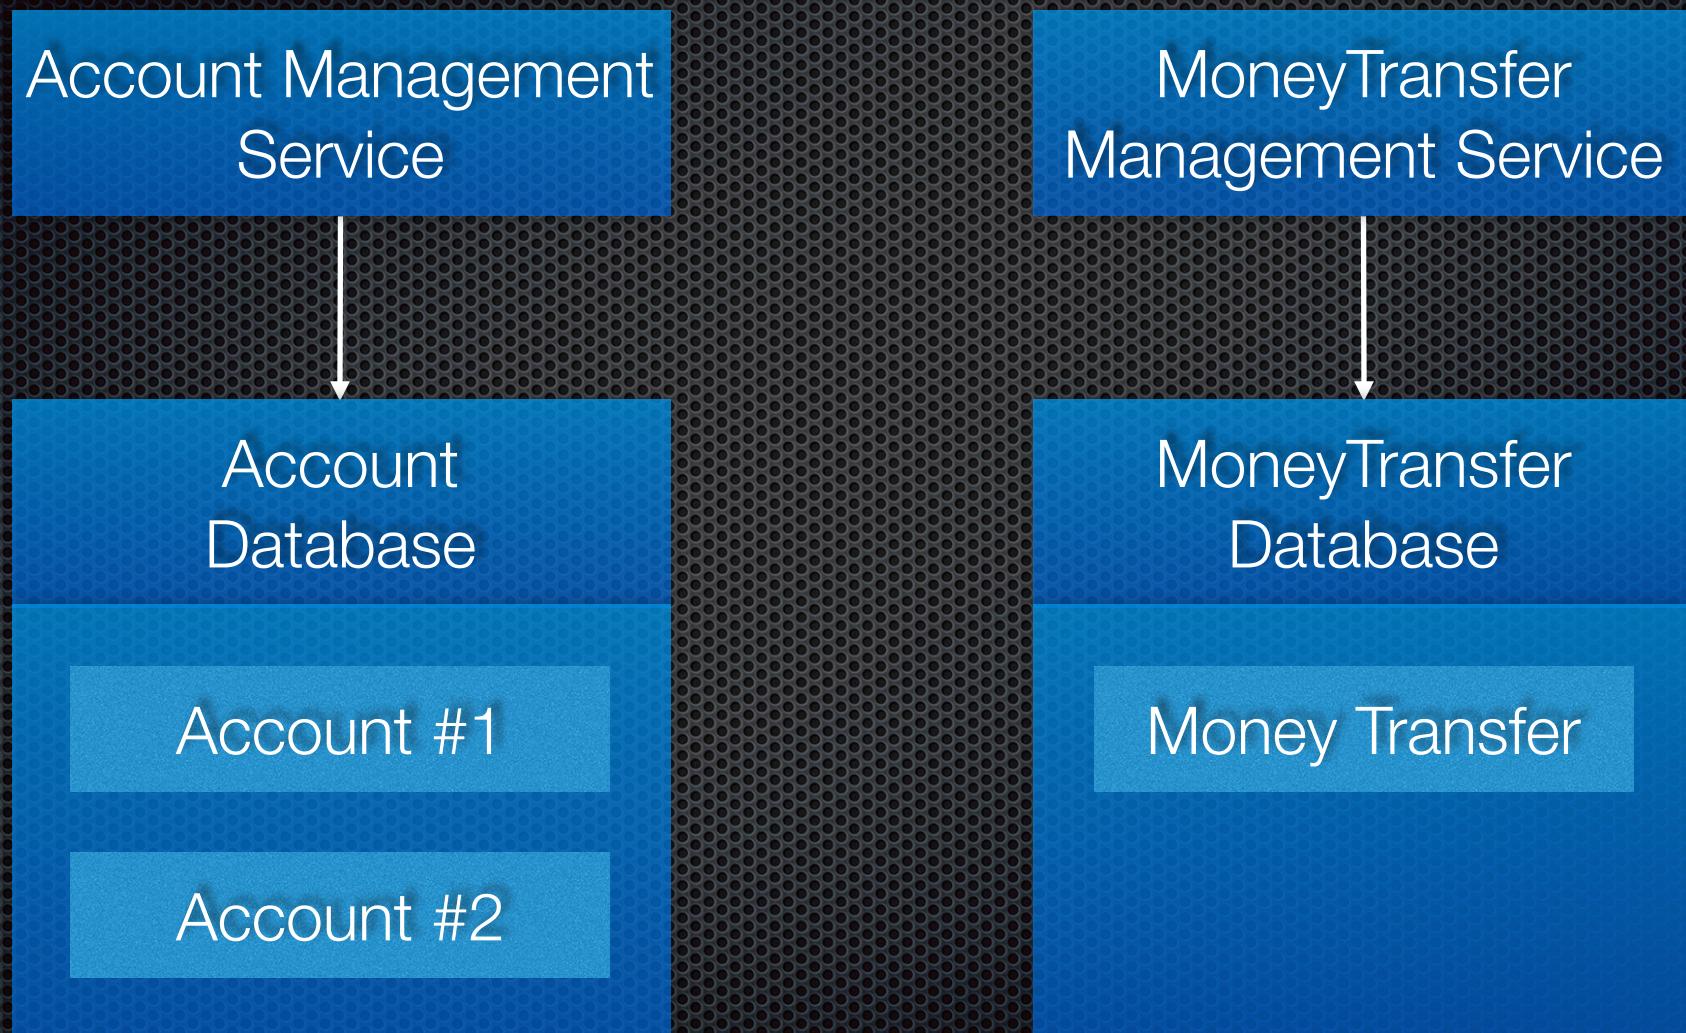
How to maintain consistency without 2PC?

Example #2 - Cassandra main table <=> index table



How to maintain consistency without 2PC?

Example #3: Money transfer



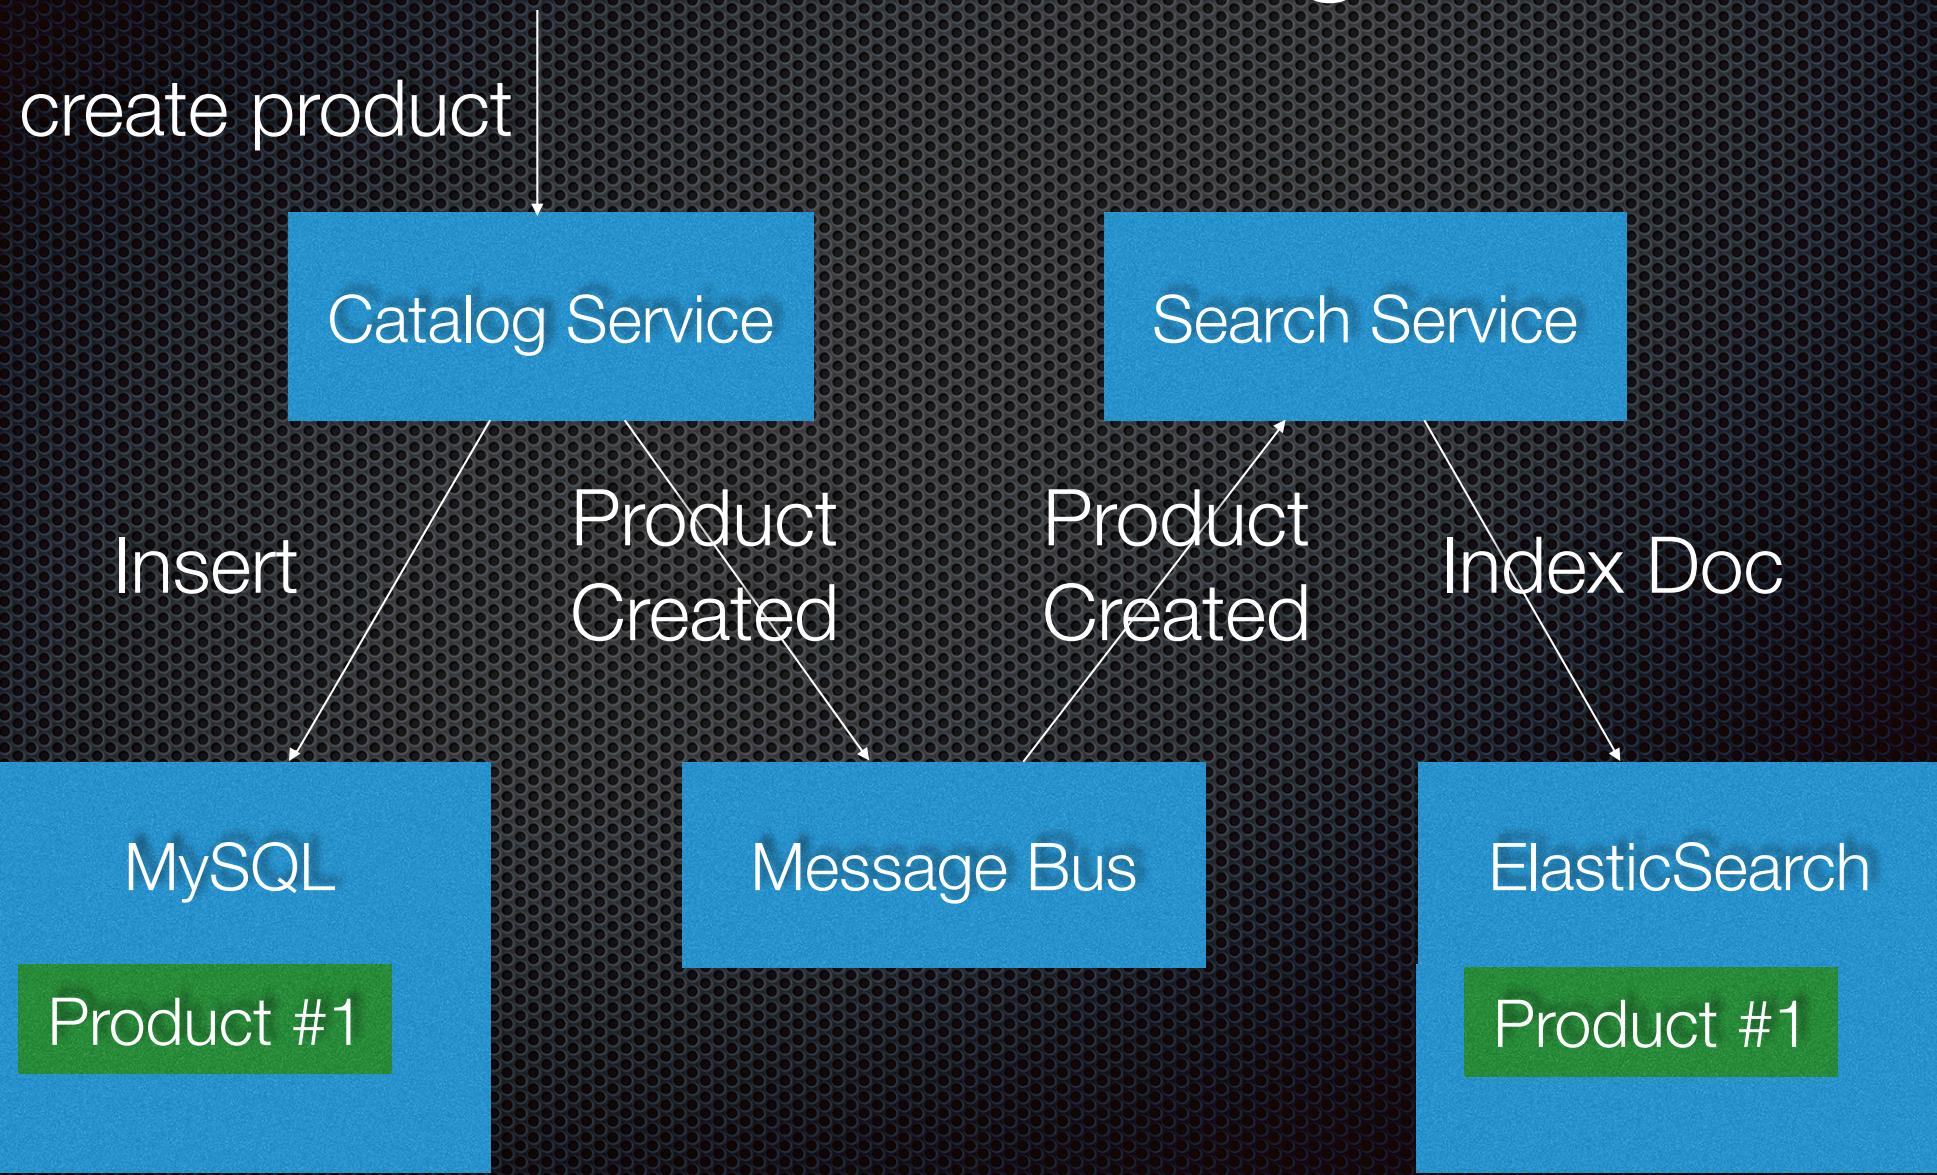
How to maintain consistency without 2PC?

@crichardson

Event-based architecture to the rescue

- Components (e.g. services) publish events when state changes
- Components subscribe to events
 - Maintains eventual consistency across multiple aggregates (in multiple datastores)
 - Synchronize replicated data

Event-driven synchronization: SQL + Text Search engine



Eventually consistent money transfer

transferMoney()



MoneyTransferService

MoneyTransfer
fromAccountId = 101
toAccountId = 202
amount = 55
state = COMPLETED

AccountService

Account
id = 101
balance = 195

Account
id = 202
balance = 180

Subscribes to:

AccountDebitedEvent
AccountCreditedEvent

publishes:

MoneyTransferCreatedEvent
DebitRecordedEvent

Subscribes to:

MoneyTransferCreatedEvent
DebitRecordedEvent

Publishes:

AccountDebitedEvent
AccountCreditedEvent

Message Bus

To maintain consistency
a service must
atomically publish an event
whenever
a domain object changes

How to atomically update the datastore and publish event(s)?

- Use 2PC
 - Guaranteed atomicity **BUT**
 - Need a distributed transaction manager
 - Database and message broker must support 2PC
 - Impacts reliability
 - Not fashionable
 - 2PC is best avoided
- Use datastore as a message queue
 1. Update database: new entity state & event
 2. Consume event & mark event as consumed
 - Eventually consistent mechanism
 - See BASE: An Acid Alternative, <http://bit.ly/ebaybase>
 - **BUT** Tangled business logic and event publishing code
 - Difficult to implement when using a NoSQL database :-(

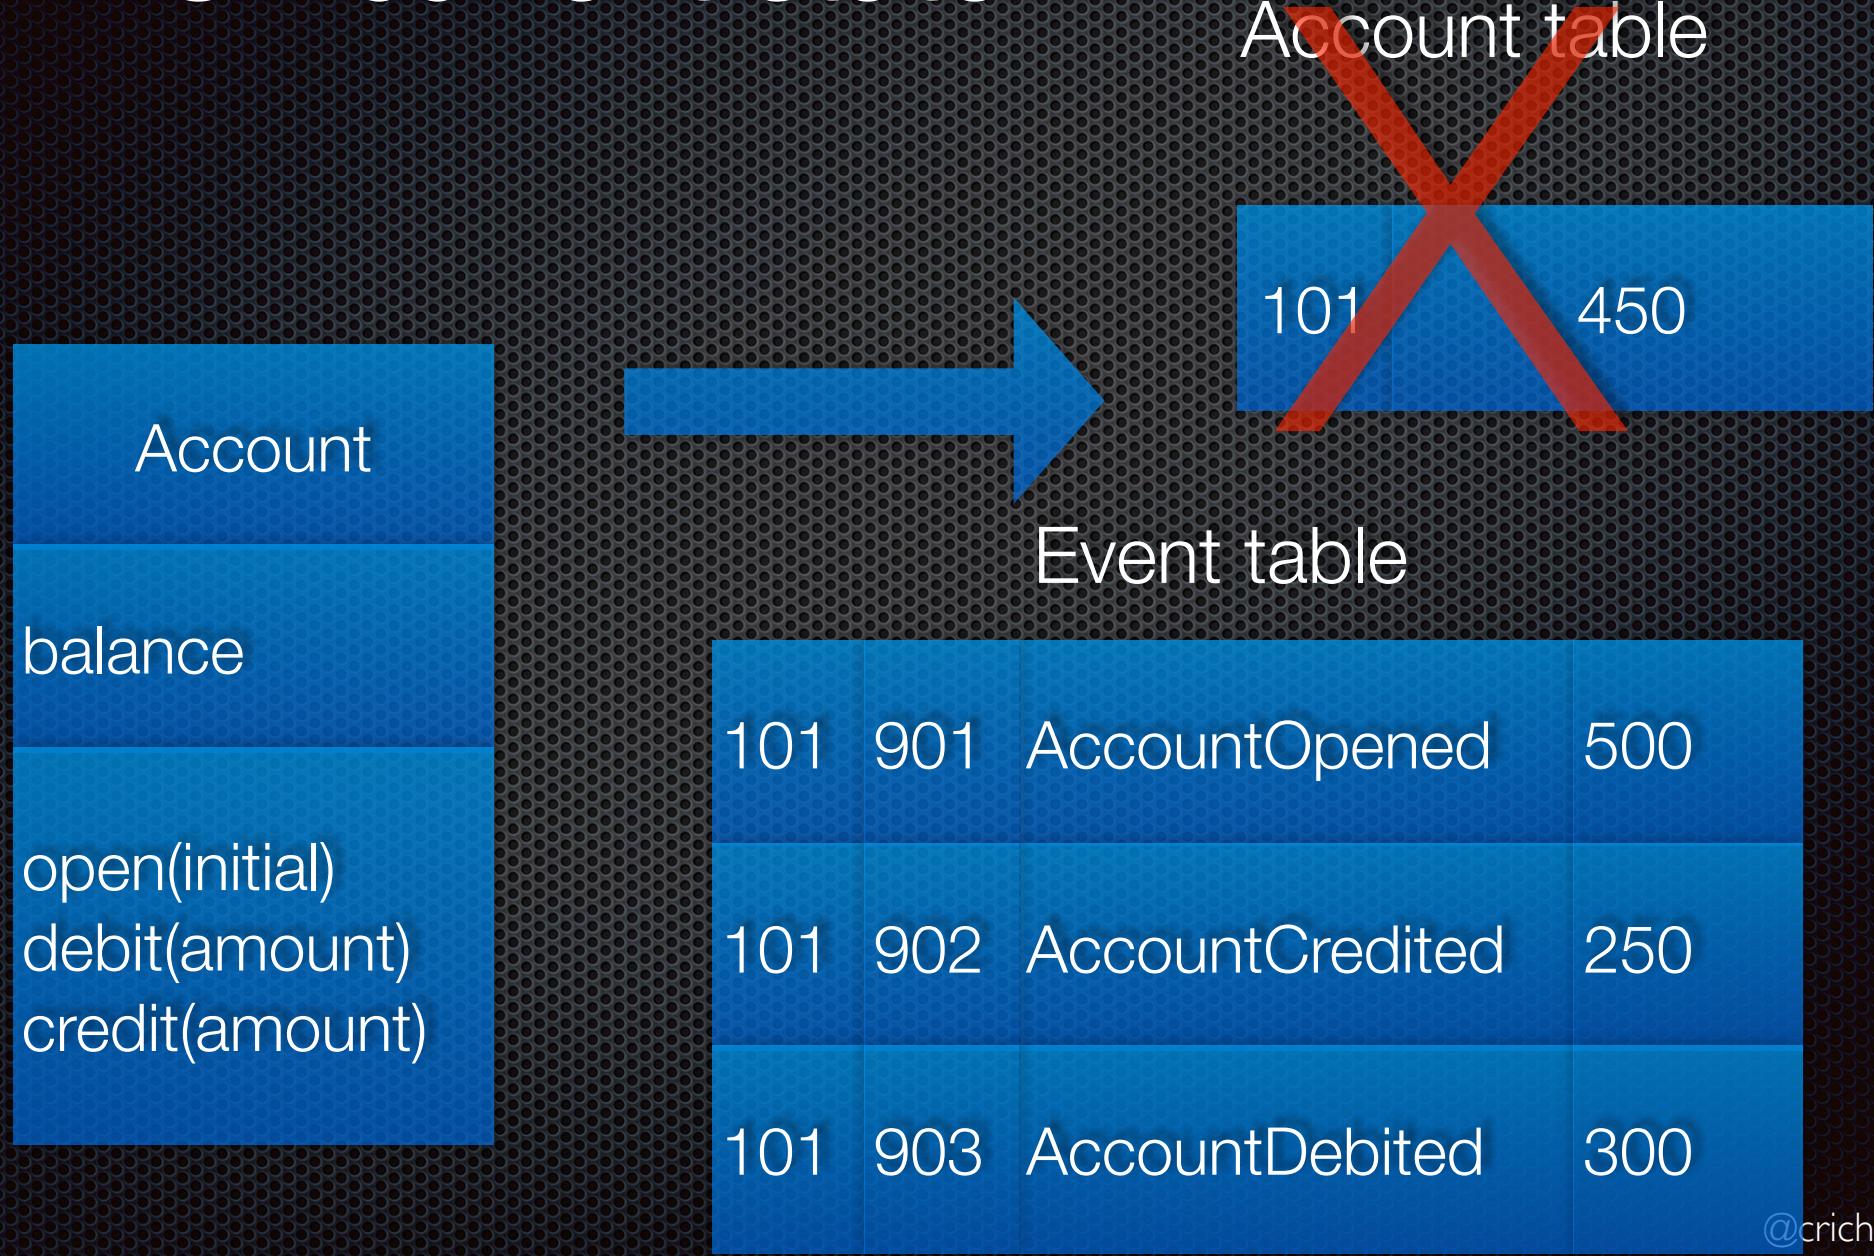
Agenda

- Why build event-driven microservices?
- Overview of event sourcing
- Designing microservices with event sourcing
- Implementing queries in an event sourced application
- Building and deploying microservices

Event sourcing

- For each aggregate:
 - Identify (state-changing) domain events
 - Define Event classes
- For example,
 - Account: AccountOpenedEvent, AccountDebitedEvent, AccountCreditedEvent
 - ShoppingCart: ItemAddedEvent, ItemRemovedEvent, OrderPlacedEvent

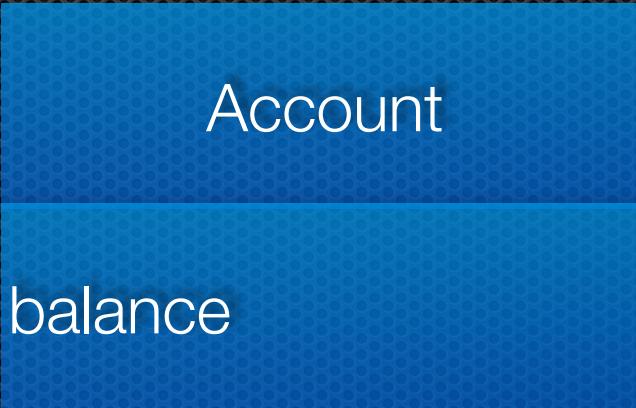
Persists events NOT current state



Replay events to recreate state

Events

AccountOpenedEvent(balance)
AccountDebitedEvent(amount)
AccountCreditedEvent(amount)



Two actions that must be atomic

Before: update state + publish events

Now: persist (and publish) events

Single action that can be done atomically

Aggregate traits

Apply event returning
updated Aggregate

```
trait Aggregate[T] { self : T =>
  def applyEvent : PartialFunction[Event, T]
}

trait CommandProcessingAggregate[T, -CT] extends Aggregate[T] { self : T =>
  def processCommand : PartialFunction[CT, Seq[Event]]
}
```

Map Command to Events

Account - command processing

```
case class Account(balance : BigDecimal)
  extends CommandProcessingAggregate[Account, AccountCommand] {

  def this() = this(null)

  import net.chrisrichardson.eventstore.examples.bank.accounts.AccountCommands._

  def processCommand = {
    case OpenedAccountCommand(initialBalance) =>
      Seq(AccountOpenedEvent(initialBalance))

    case CreditAccountCommand(amount, transactionId) =>
      Seq(AccountCreditedEvent(amount, transactionId))

    case DebitAccountCommand(amount, transactionId) if amount <= balance =>
      Seq(AccountDebitedEvent(amount, transactionId))

    case DebitAccountCommand(amount, transactionId) =>
      Seq(AccountDebitFailedDueToInsufficientFundsEvent(amount, transactionId))
  }
}
```

Prevent
overdraft

Account - applying events

Immutable

```
case class Account(balance : BigDecimal)
  extends CommandProcessingAggregate[Account, AccountCommand] {

  def applyEvent = {

    case AccountOpenedEvent(initialBalance) => copy(balance = initialBalance)

    case AccountDebitedEvent(amount, _) => copy(balance = balance - amount)

    case AccountCreditedEvent(amount, _) =>
      copy(balance = balance + amount)

    case AccountDebitFailedDueToInsufficientFundsEvent(amount, _) =>
      this
  }
}
```

Request handling in an event-sourced application

Microservice A

`pastEvents = findEvents(entityId)`

`new()`

`applyEvents(pastEvents)`

`newEvents = processCmd(SomeCmd)`

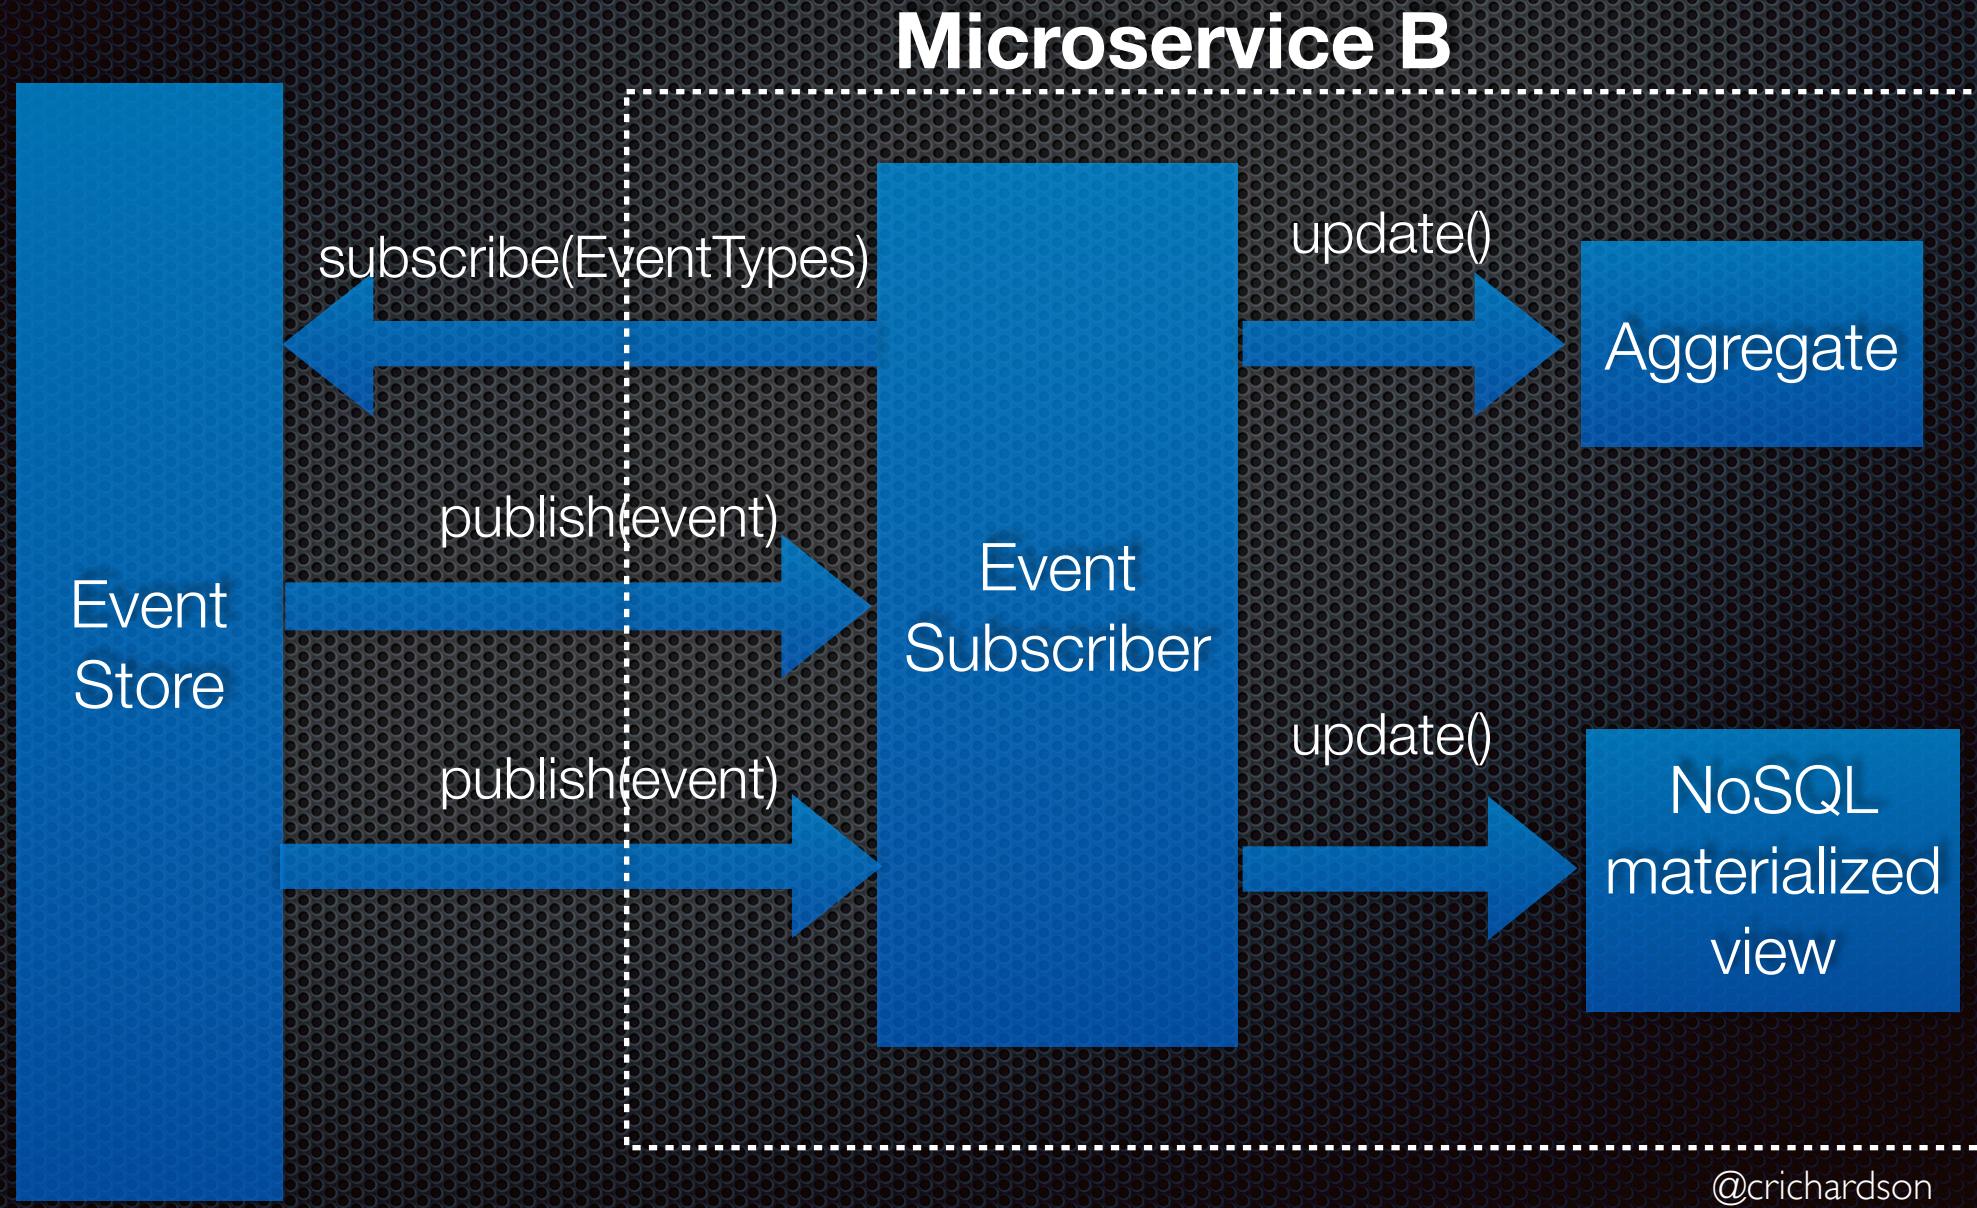
`saveEvents(newEvents)`

HTTP
Handler

Account

Event
Store

Event Store publishes events - consumed by other services



Optimizing using snapshots

- Most aggregates have relatively few events
- BUT consider a 10-year old Account ⇒ many transactions
- Therefore, use snapshots:
 - Periodically save snapshot of aggregate state
 - Typically serialize a memento of the aggregate
 - Load latest snapshot + subsequent events

Event Store API

```
trait EventStore {  
  
    def save[T <: Aggregate[T]](entity: T, events: Seq[Event],  
        assignedId : Option[EntityId] = None): Future[EntityWithIdAndVersion[T]]  
  
    def update[T <: Aggregate[T]](entityIdAndVersion : EntityIdAndVersion,  
        entity: T, events: Seq[Event]): Future[EntityWithIdAndVersion[T]]  
  
    def find[T <: Aggregate[T] : ClassTag](entityId: EntityId) :  
        Future[EntityWithIdAndVersion[T]]  
  
    def findOptional[T <: Aggregate[T] : ClassTag](entityId: EntityId)  
        Future[Option[EntityWithIdAndVersion[T]]]  
  
    def subscribe(subscriptionId: SubscriptionId):  
        Future[AcknowledgableEventStream]  
}
```

Event store implementations

- Home-grown/DIY
- geteventstore.com by Greg Young
- Talk to me about my project :-)

Business benefits of event sourcing

- Built-in, reliable audit log
- Enables temporal queries
- Publishes events needed by big data/predictive analytics etc.
- Preserved history ⇒ More easily implement future requirements

Technical benefits of event sourcing

- Solves data consistency issues in a Microservice/NoSQL-based architecture:
 - Atomically save and publish events
 - Event subscribers update other aggregates ensuring eventual consistency
 - Event subscribers update materialized views in SQL and NoSQL databases (more on that later)
- Eliminates O/R mapping problem

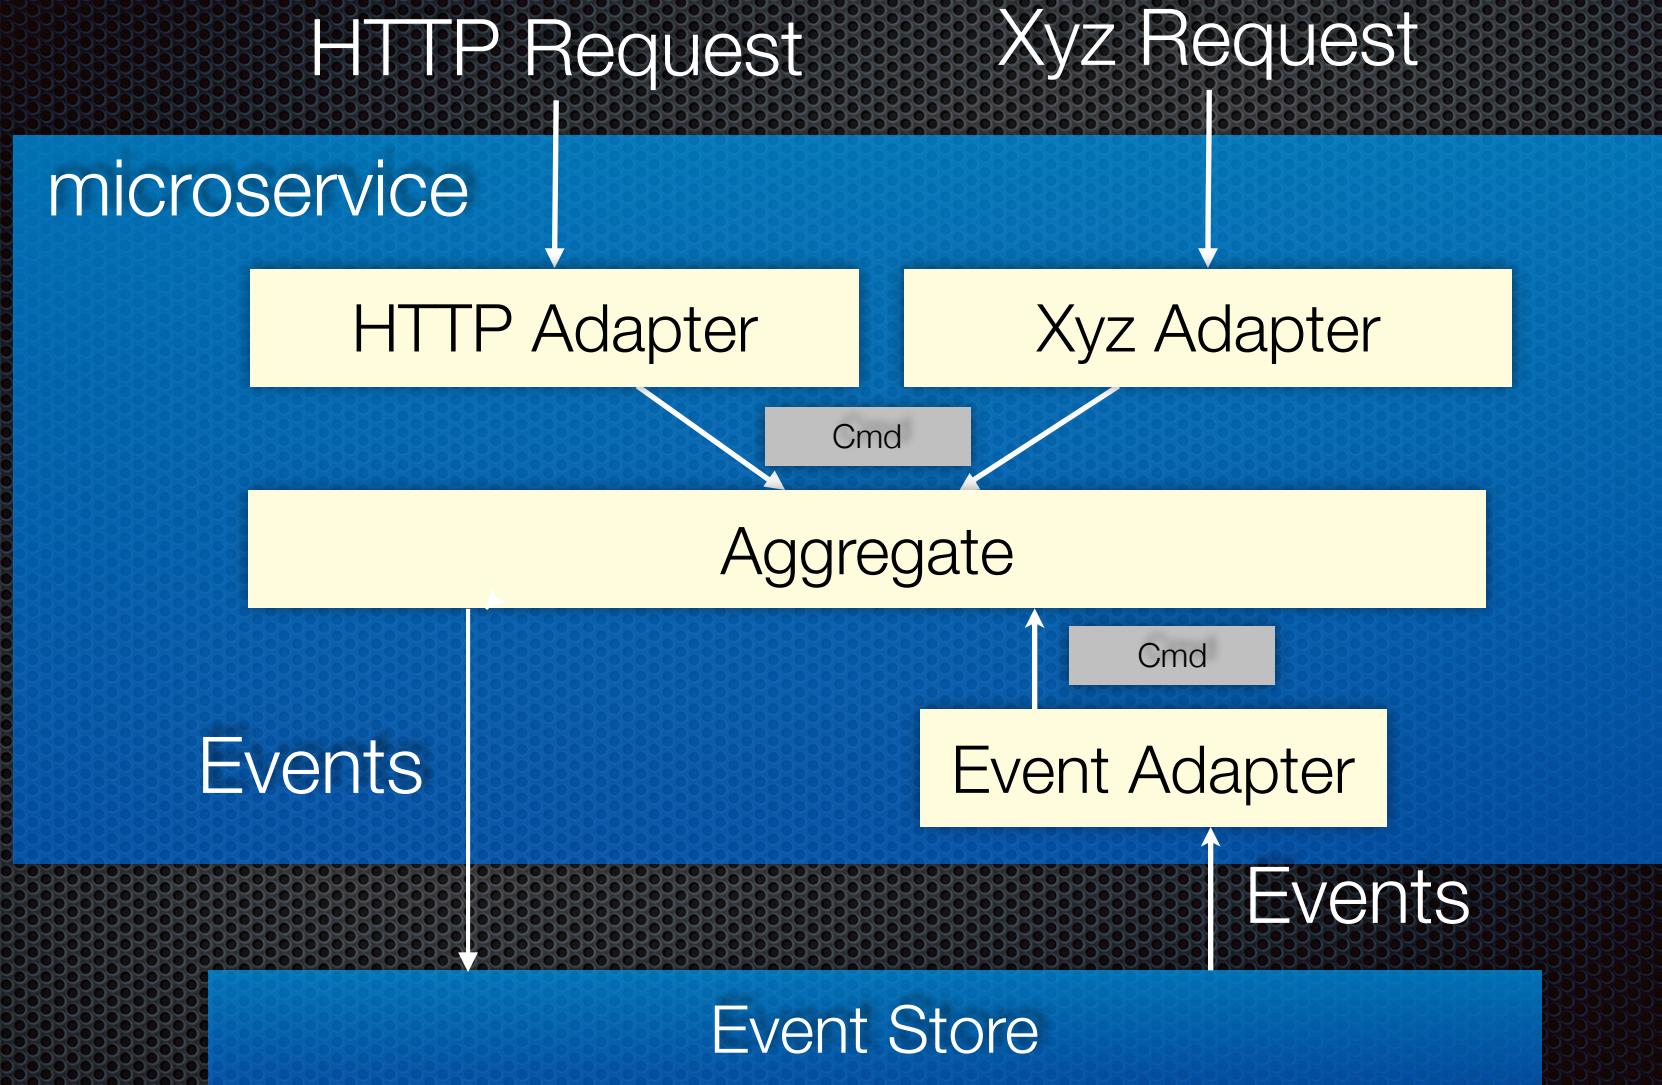
Drawbacks of event sourcing

- ❖ Weird and unfamiliar
- ❖ Events = a historical record of your bad design decisions
- ❖ Handling duplicate events can be tricky
- ❖ Application must handle eventually consistent data
- ❖ Event store only directly supports PK-based lookup (more on that later)

Agenda

- Why build event-driven microservices?
- Overview of event sourcing
- Designing microservices with event sourcing
- Implementing queries in an event sourced application
- Building and deploying microservices

The anatomy of a microservice



Asynchronous Spring MVC controller

```
@RestController
class MoneyTransferController @Autowired()(moneyTransferService : MoneyTransferService,
                                             eventStore : EventStore) {

    @RequestMapping(value=Array("/transfers"), method = Array(RequestMethod.POST))
    def create(@RequestBody transferDetails : TransferDetails) = WebUtil.toDeferredResult {
        for (transaction <- moneyTransferService.transferMoney(transferDetails))
            yield CreateMoneyTransferResponse(transaction.entityId.id)
    }
}
```

MoneyTransferService

```
class MoneyTransferService(implicit eventStore : EventStore) {  
    def transferMoney(transferDetails : TransferDetails) =  
        newEntity[MoneyTransfer] <== CreateMoneyTransferCommand(transferDetails)  
}
```

DSL concisely specifies:

1. Creates MoneyTransfer aggregate
2. Processes command
3. Applies events
4. Persists events

MoneyTransfer Aggregate

```
case class TransferDetails(fromAccountId : EntityId, toAccountId : EntityId, amount : BigDecimal)

case class MoneyTransfer(state : TransferStates.State, details : TransferDetails)
  extends CommandProcessingAggregate[MoneyTransfer, MoneyTransferCommand] {

  def this() = this(TransferStates.NEW, null)

  import net.chrisrichardson.eventstore.examples.bank.transactions.MoneyTransferCommands._

  def processCommand = {...}

  def applyEvent = {...}

}
```

```
case class MoneyTransferCreatedEvent(details : TransferDetails) extends Event
case class DebitRecordedEvent(details : TransferDetails) extends Event
case class CreditRecordedEvent(details : TransferDetails) extends Event
case class TransferFailedDueToInsufficientFundsEvent() extends Event
```

Handling events published by Accounts

```
class MoneyTransferEventHandlers(implicit eventStore: EventStore)
  extends CompoundEventHandler {

  val recordDebit =
    handlerForEvent[AccountDebitedEvent] { de =>
      existingEntity[MoneyTransfer](de.event.transactionId) <==  

        RecordDebitCommand(de.entityId)
    }
}
```

- 1.Load MoneyTransfer aggregate
- 2.Processes command
- 3.Applies events
- 4.Persists events

Agenda

- Why build event-driven microservices?
- Overview of event sourcing
- Designing microservices with event sourcing
- Implementing queries in an event sourced application
- Building and deploying microservices

Let's imagine that you want to display an account and its recent transactions...

Displaying balance + recent credits and debits

- We need to do a “join: between the Account and the corresponding MoneyTransfers
- (Assuming Debit/Credit events don’t include other account, ...)

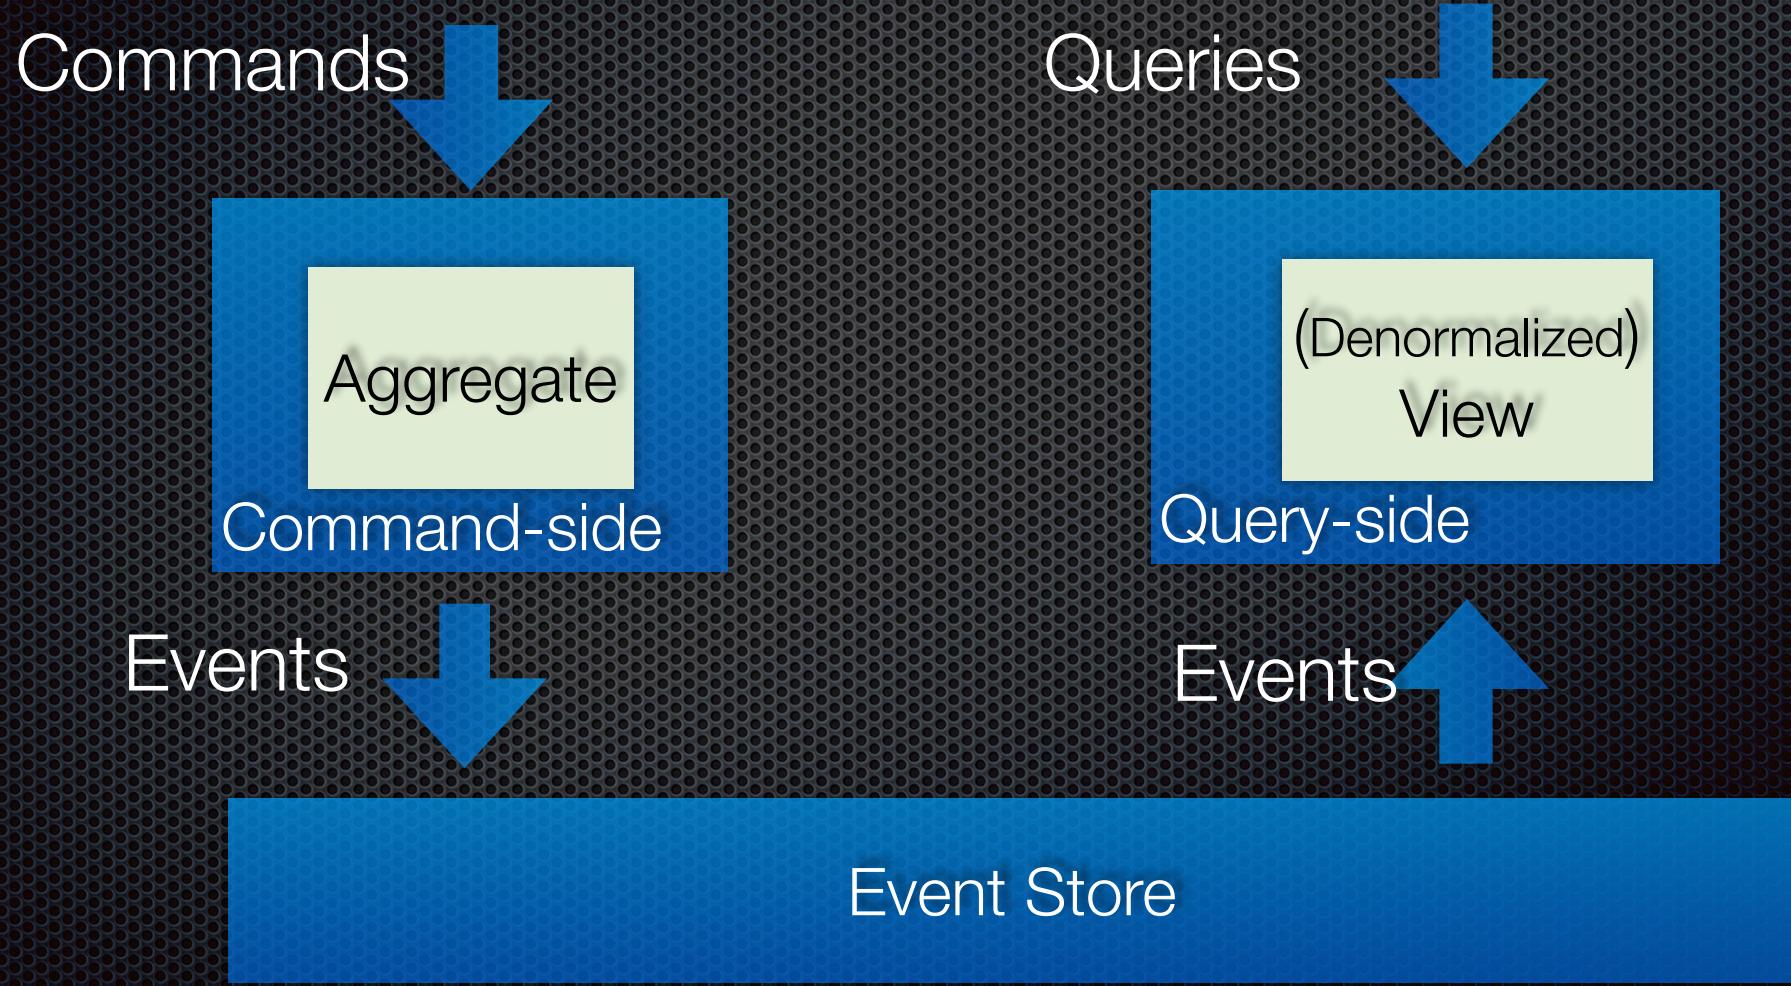
BUT

- Event Store = primary key lookup of individual aggregates, ...

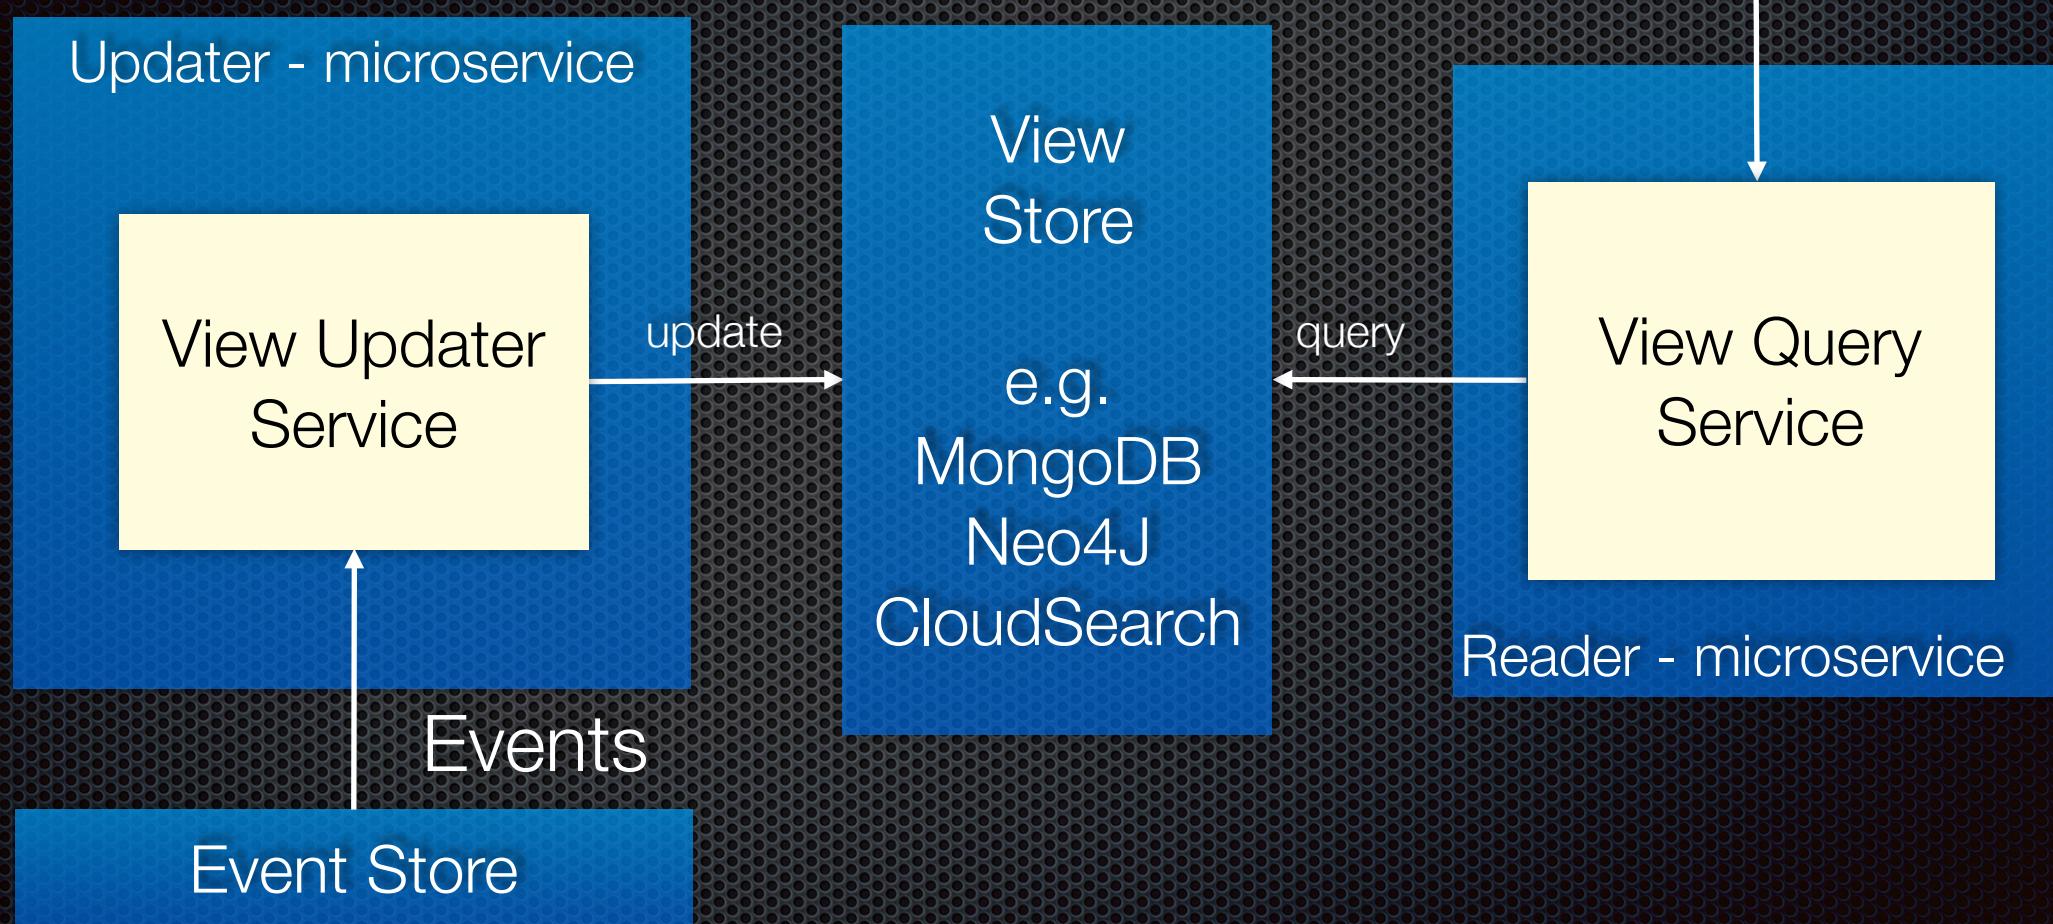


- Use **C**ommand **Q**uery **R**esponsibility **S**egregation

Command Query Responsibility Segregation (CQRS)



Query-side microservices



Persisting account balance and recent transactions in MongoDB

```
{  
  id: "298993498",  
  balance: 100000,  
  transfers : [  
    {"transferId" : "4552840948484",  
     "fromAccountId" : 298993498,  
     "toAccountId" : 3483948934,  
     "amount" : 5000}, ...  
  ],  
  changes: [  
    {"changeId" : "93843948934",  
     "transferId" : "4552840948484",  
     "transactionType" : "AccountDebited",  
     "amount" : 5000}, ...  
  ]  
}
```

Current balance

MoneyTransfers that update the account

The debits and credits

Denormalized = efficient Lookup

@crichardson

Persisting account info using MongoDB...

```
class AccountInfoUpdateService
  (accountInfoRepository : AccountInfoRepository, mongoTemplate : MongoTemplate)
    extends CompoundEventHandler {

  @EventHandlerMethod
  def created(de: DispatchedEvent[AccountOpenedEvent]) = ...

  @EventHandlerMethod
  def recordDebit(de: DispatchedEvent[AccountDebitedEvent]) = ...

  @EventHandlerMethod
  def recordCredit(de: DispatchedEvent[AccountCreditedEvent]) = ...

  @EventHandlerMethod
  def recordTransfer(de: DispatchedEvent[MoneyTransferCreatedEvent]) = ...

}
```

Other kinds of views

- AWS Cloud Search
 - Text search as-a-Service
 - View updaters batches aggregates to index
 - View query service does text search
- AWS DynamoDB
 - NoSQL as-a-Service
 - On-demand scalable - specify desired read/write capacity
 - Document and key-value data models
 - Useful for denormalized, UI oriented views

Benefits and drawbacks of CQRS

Benefits

- Necessary in an event-sourced architecture
- Separation of concerns = simpler command and query models
- Supports multiple denormalized views
- Improved scalability and performance

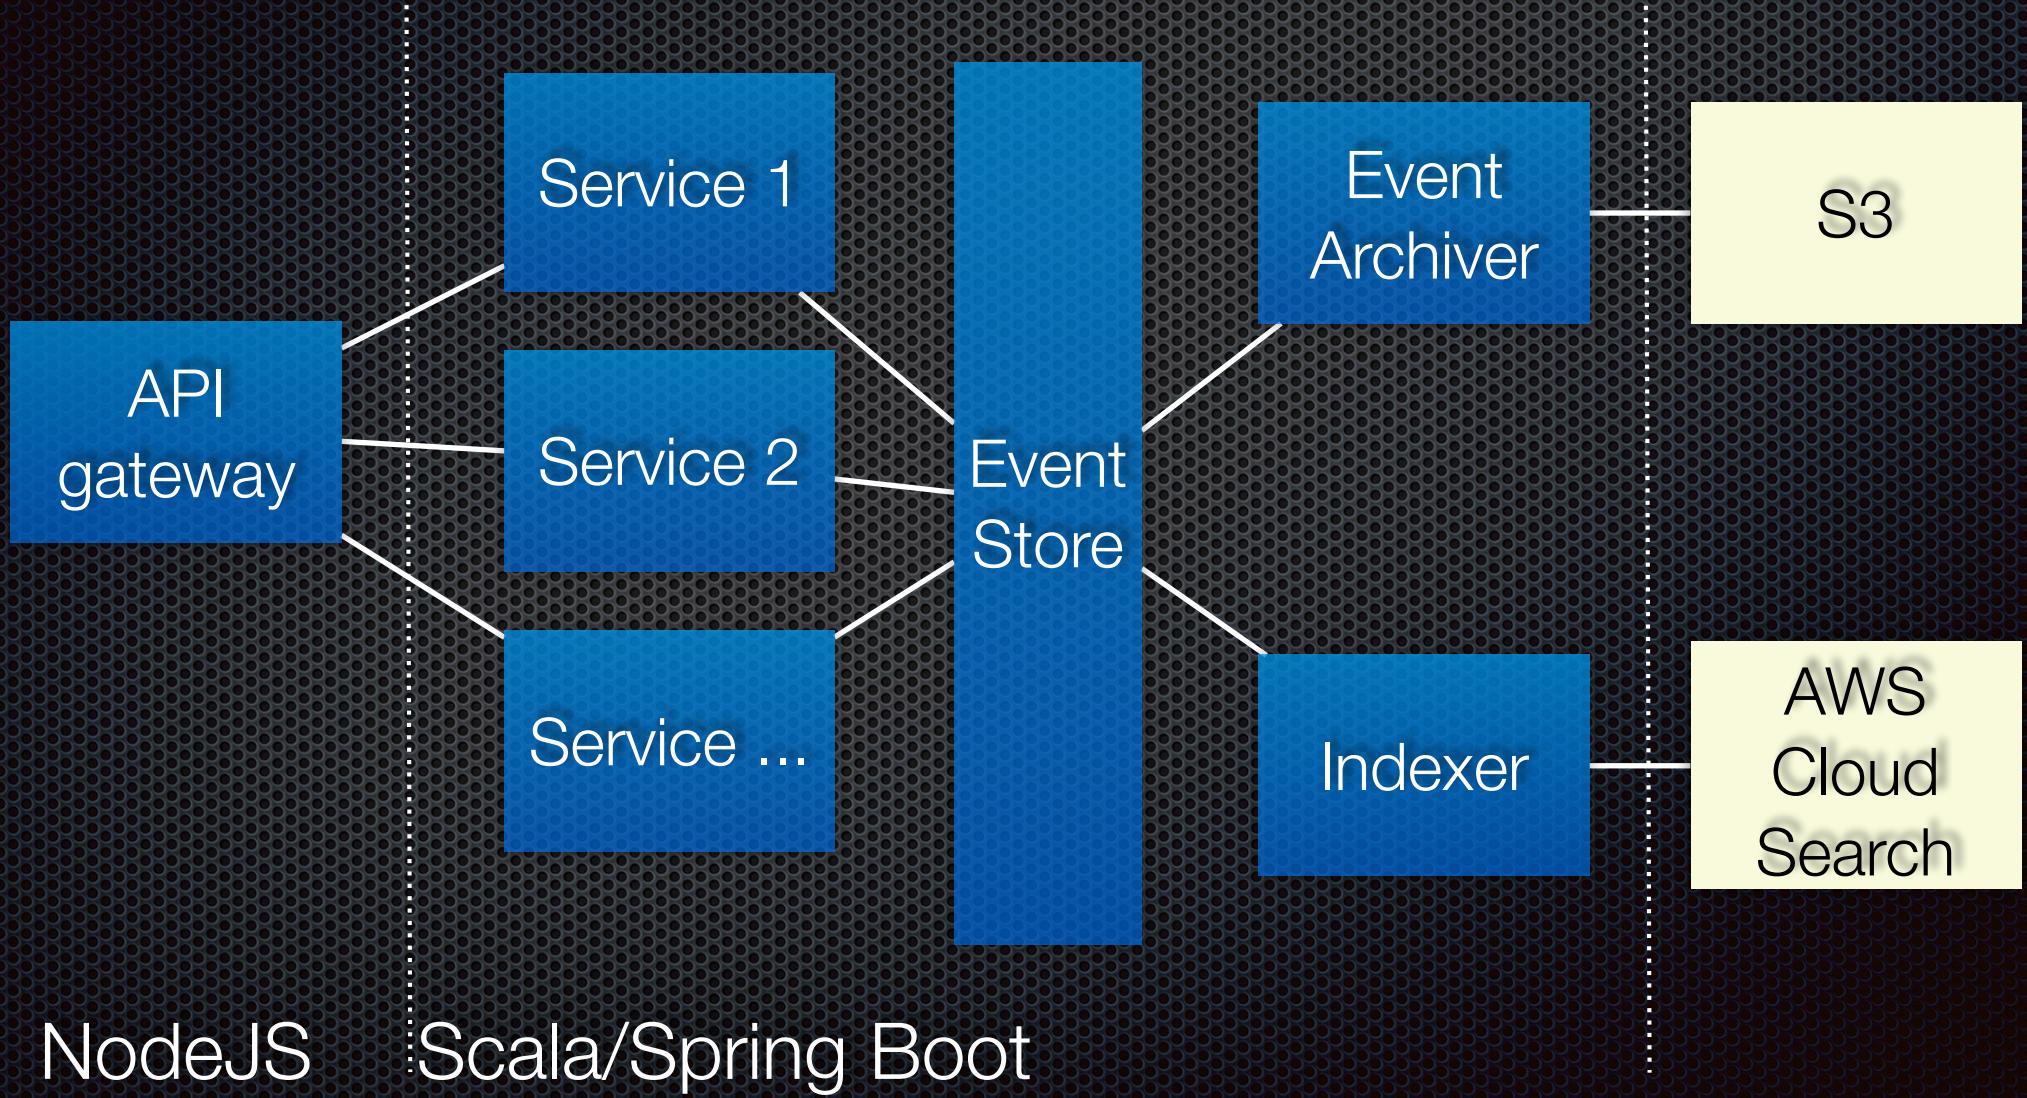
Drawbacks

- Complexity
- Potential code duplication
- Replication lag/eventually consistent views

Agenda

- Why build event-driven microservices?
- Overview of event sourcing
- Designing microservices with event sourcing
- Implementing queries in an event sourced application
- Building and deploying microservices

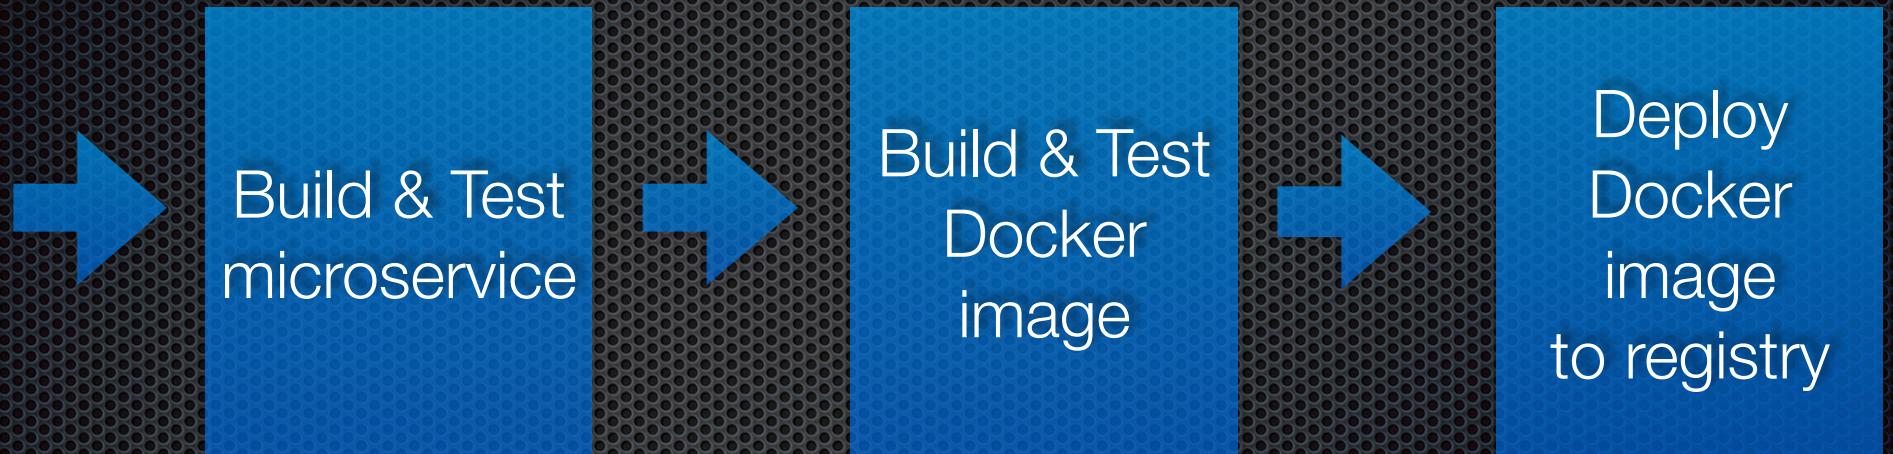
My application architecture



Spring Boot based micro services

- Makes it easy to create stand-alone, production ready microservices
- Automatically configures Spring using Convention over Configuration
- Externalizes configuration
- Built-in health checks and (Codahale) metrics
- Generates standalone executable JARs with embedded web server

Jenkins-based deployment pipeline



One pipeline per microservice

Building Docker images

docker/build.sh

```
cp ../build/libs/service.${1}.jar build/service.jar
```

```
docker build -t service-${VERSION} .
```

*Building only takes 5
seconds!*

Smoke testing docker images



Docker daemon must listen on
TCP port

Publishing Docker images

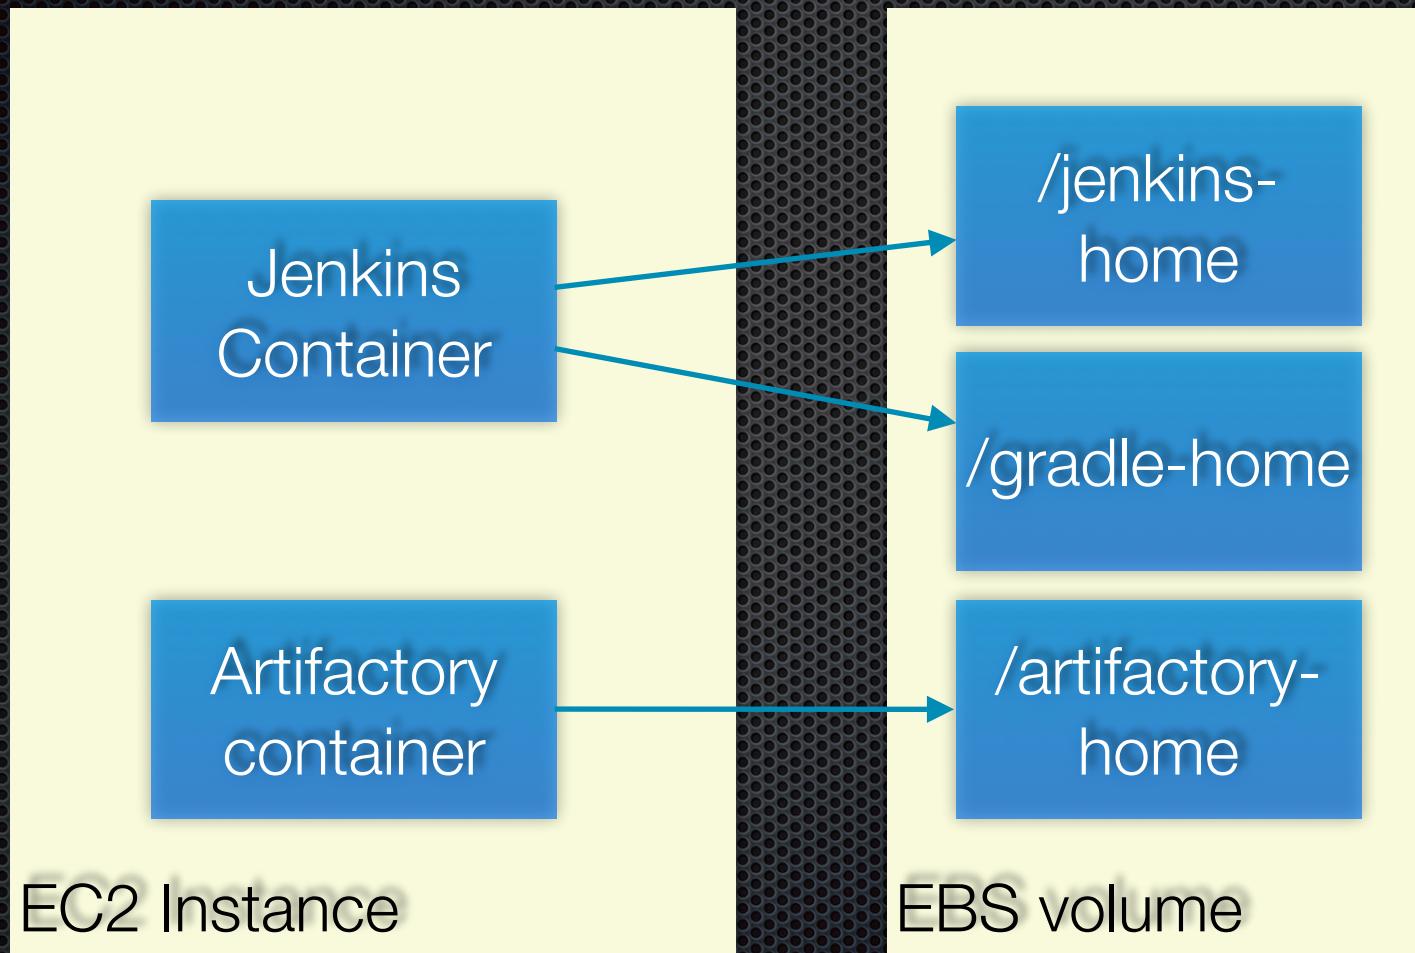
docker/publish.sh

```
docker tag service-${VERSION}:latest \  
 ${REGISTRY_HOST_AND_PORT}/service-${VERSION}
```

```
docker push ${REGISTRY_HOST_AND_PORT}/service-${VERSION}
```

Pushing only takes 25
seconds!

CI environment runs on Docker



Updating the production environment

- Large EC2 instance running Docker
- Deployment tool:
 1. Compares running containers with what's been built by Jenkins
 2. Pulls latest images from Docker registry
 3. Stops old versions
 4. Launches new versions
- One day: use Docker clustering solution and a service discovery mechanism,
 - Most likely, AWS container service
 - Mesos and Marathon + Zookeeper, Kubernetes or ???

Summary

- Event sourcing solves key data consistency issues with:
 - Microservices
 - Partitioned SQL/NoSQL databases
- Use CQRS to implement materialized views for queries
- Spring Boot is a great foundation for microservices
- Docker is a great way to package microservices

• @crichton chris@chrisrichardson.net



<http://plainoldobjects.com>

<http://microservices.io>